


# **Chapter 22**

## **Concurrency Control Techniques**



Sixth Edition

# Fundamentals of Database Systems

Elmasri • Navathe

Addison-Wesley  
is an imprint of

PEARSON

Copyright © 2011 Pearson Education, Inc. Publishing as Pearson Addison-Wesley

# Database Concurrency Control

- 1 Purpose of Concurrency Control
  - To enforce Isolation (through mutual exclusion) among conflicting transactions.
  - To preserve database consistency through consistency preserving execution of transactions.
  - To resolve read-write and write-write conflicts.
- Example:
  - In concurrent execution environment if T1 conflicts with T2 over a data item A, then the existing concurrency control decides if T1 or T2 should get the A and if the other transaction is rolled-back or waits.

# Database Concurrency Control

## Two-Phase Locking Techniques

- Locking is an operation which secures
  - (a) permission to Read
  - (b) permission to Write a data item for a transaction.
- Example:
  - Lock (X). Data item X is locked in behalf of the requesting transaction.
- Unlocking is an operation which removes these permissions from the data item.
- Example:
  - Unlock (X): Data item X is made available to all other transactions.
- Lock and Unlock are Atomic operations.

# Database Concurrency Control

## Two-Phase Locking Techniques: Essential components

- **Two locks modes:**
  - (a) shared (read)                      (b) exclusive (write).
- **Shared mode: shared lock (X)**
  - More than one transaction can apply share lock on X for reading its value but no write lock can be applied on X by any other transaction.
- **Exclusive mode: Write lock (X)**
  - Only one write lock on X can exist at any time and no shared lock can be applied by any other transaction on X.
- **Conflict matrix**

	Read	Write
Read	Y	N
Write	N	N

# Database Concurrency Control

## Two-Phase Locking Techniques: Essential components

- **Lock Manager:**
  - Managing locks on data items.
- **Lock table:**
  - Lock manager uses it to store the identify of transaction locking a data item, the data item, lock mode and pointer to the next data item locked. One simple way to implement a lock table is through linked list.

Transaction ID	Data item id	lock mode	Ptr to next data item
T1	X1	Read	Next

# Database Concurrency Control

## Two-Phase Locking Techniques: Essential components

- Database requires that all transactions should be well-formed. A transaction is well-formed if:
  - It must lock the data item before it reads or writes to it.
  - It must not lock an already locked data items and it must not try to unlock a free data item.

# Database Concurrency Control

## Two-Phase Locking Techniques: Essential components

- The following code performs the lock operation:

```
B: if LOCK (X) = 0 (*item is unlocked*)  
    then LOCK (X) ← 1 (*lock the item*)  
    else begin  
        wait (until lock (X) = 0 and  
            the lock manager wakes up the transaction);  
    goto B  
end;
```

# Database Concurrency Control

## Two-Phase Locking Techniques: Essential components

- The following code performs the unlock operation:

$\text{LOCK}(X) \leftarrow 0$  (\*unlock the item\*)

if any transactions are waiting then

wake up one of the waiting the transactions;



# Database Concurrency Control

## Two-Phase Locking Techniques: Essential components

- The following code performs the read operation:

```
B: if LOCK (X) = "unlocked" then
  begin LOCK (X) ← "read-locked";
    no_of_reads (X) ← 1;
  end
else if LOCK (X) = "read-locked" then
  no_of_reads (X) ← no_of_reads (X) + 1
  else begin wait (until LOCK (X) = "unlocked" and
    the lock manager wakes up the transaction);
    go to B
  end;
```

# Database Concurrency Control

## Two-Phase Locking Techniques: Essential components

- The following code performs the write lock operation:

```
B: if LOCK(X) = "unlocked"
    then LOCK(X) ← "write-locked"
else begin wait (until LOCK(X) = "unlocked" and
    the lock manager wakes up the transaction);
    go to B
end;
```

# Database Concurrency Control

## Two-Phase Locking Techniques: Essential components

- The following code performs the unlock operation:

```
if LOCK (X) = "write-locked" then
begin LOCK (X) ← "unlocked";
    wakes up one of the transactions, if any
end
else if LOCK (X) ← "read-locked" then
begin
    no_of_reads (X) ← no_of_reads (X) -1
    if no_of_reads (X) = 0 then
begin
    LOCK (X) = "unlocked";
    wake up one of the transactions, if any
end
end;
end;
```

# Database Concurrency Control

## Two-Phase Locking Techniques: Essential components

- Lock conversion
  - Lock upgrade: existing read lock to write lock
    - if  $T_i$  has a read-lock (X) and  $T_j$  has no read-lock (X) ( $i \neq j$ ) then  
convert read-lock (X) to write-lock (X)
    - else  
force  $T_i$  to wait until  $T_j$  unlocks X
  - Lock downgrade: existing write lock to read lock
    - $T_i$  has a write-lock (X) (\*no transaction can have any lock on X\*)
    - convert write-lock (X) to read-lock (X)

# Database Concurrency Control

## Two-Phase Locking Techniques: The algorithm

- Two Phases:
  - (a) Locking (Growing)
  - (b) Unlocking (Shrinking).
- **Locking (Growing) Phase:**
  - A transaction applies locks (read or write) on desired data items one at a time.
- **Unlocking (Shrinking) Phase:**
  - A transaction unlocks its locked data items one at a time.
- **Requirement:**
  - For a transaction these two phases must be mutually exclusively, that is, during locking phase unlocking phase must not start and during unlocking phase locking phase must not begin.

# Database Concurrency Control

## Two-Phase Locking Techniques: The algorithm

### T1

read\_lock (Y);  
read\_item (Y);  
unlock (Y);  
write\_lock (X);  
read\_item (X);  
X:=X+Y;  
write\_item (X);  
unlock (X);

### T2

read\_lock (X);  
read\_item (X);  
unlock (X);  
write\_lock (Y);  
read\_item (Y);  
Y:=X+Y;  
write\_item (Y);  
unlock (Y);

### Result

Initial values: X=20; Y=30  
Result of serial execution  
T1 followed by T2  
X=50, Y=80.  
Result of serial execution  
T2 followed by T1  
X=70, Y=50

# Database Concurrency Control

## Two-Phase Locking Techniques: The algorithm

T1	T2	<u>Result</u>
<code>read_lock (Y);</code> <code>read_item (Y);</code> <code>unlock (Y);</code>	<code>read_lock (X);</code> <code>read_item (X);</code> <code>unlock (X);</code> <code>write_lock (Y);</code> <code>read_item (Y);</code> <code>Y:=X+Y;</code> <code>write_item (Y);</code> <code>unlock (Y);</code>	$X=50; Y=50$ Nonserializable because it. violated two-phase policy.
<div data-bbox="92 748 195 786">Time</div> <div data-bbox="208 665 227 951"></div>		
<code>write_lock (X);</code> <code>read_item (X);</code> <code>X:=X+Y;</code> <code>write_item (X);</code> <code>unlock (X);</code>		

# Database Concurrency Control

## Two-Phase Locking Techniques: The algorithm

### T'1

```
read_lock (Y);  
read_item (Y);  
write_lock (X);  
unlock (Y);  
read_item (X);  
X:=X+Y;  
write_item (X);  
unlock (X);
```

### T'2

```
read_lock (X);  
read_item (X);  
Write_lock (Y);  
unlock (X);  
read_item (Y);  
Y:=X+Y;  
write_item (Y);  
unlock (Y);
```

T1 and T2 follow two-phase policy but they are subject to deadlock, which must be dealt with.



# Database Concurrency Control

## Two-Phase Locking Techniques: The algorithm

- Two-phase policy generates two locking algorithms
  - (a) **Basic**
  - (b) **Conservative**
- **Conservative:**
  - Prevents deadlock by locking all desired data items before transaction begins execution.
- **Basic:**
  - Transaction locks data items incrementally. This may cause deadlock which is dealt with.
- **Strict:**
  - A more stricter version of Basic algorithm where unlocking is performed after a transaction terminates (commits or aborts and rolled-back). This is the most commonly used two-phase locking algorithm.

# Database Concurrency Control

## Dealing with Deadlock and Starvation

- **Deadlock**

T'1

read\_lock (Y);  
read\_item (Y);

write\_lock (X);  
(waits for X)

T'2

read\_lock (X);  
read\_item (X);

write\_lock (Y);  
(waits for Y)

T1 and T2 did follow two-phase policy but they are deadlock

- **Deadlock (T'1 and T'2)**

# Database Concurrency Control

## Dealing with Deadlock and Starvation

### ■ **Deadlock prevention**

- A transaction locks all data items it refers to before it begins execution.
- This way of locking prevents deadlock since a transaction never waits for a data item.
- The conservative two-phase locking uses this approach.

# Database Concurrency Control


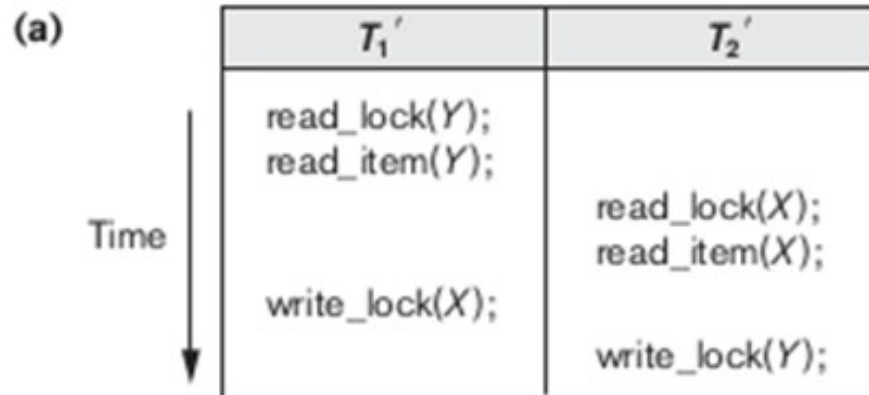
## Dealing with Deadlock and Starvation

### Deadlock detection and resolution

In this approach, deadlocks are allowed to happen. The scheduler maintains a wait-for-graph for detecting cycle. If a cycle exists, then one transaction involved in the cycle is selected (victim) and rolled-back.

To construct and maintain a wait-for graph.

- One node is created in the wait-for graph for each transaction that is currently executing.
- Whenever a transaction  $T_i$  is waiting to lock an item  $X$  that is currently locked by a transaction  $T_j$ , a directed edge ( $T_i \rightarrow T_j$ ) is created in the wait-for graph.
- When  $T_j$  releases the lock(s) on the items that  $T_i$  was waiting for, the directed edge is dropped from the wait-for graph.
- We have a state of deadlock if and only if the wait-for graph has a cycle.



Illustrating the deadlock problem. (a) A partial schedule of  $T_1'$  and  $T_2'$  that is in a state of deadlock. (b) A wait-for graph for the partial schedule in (a).

# Database Concurrency Control

## Dealing with Deadlock and Starvation

### ■ **Deadlock avoidance**

- There are many variations of two-phase locking algorithm.
- Some avoid deadlock by not letting the cycle to complete.
- That is as soon as the algorithm discovers that blocking a transaction is likely to create a cycle, it rolls back the transaction.
- Wound-Wait and Wait-Die algorithms use timestamps to avoid deadlocks by rolling-back victim.

# Database Concurrency Control

## Dealing with Deadlock and Starvation

### **Deadlock avoidance**

- Suppose that transaction  $T_i$  tries to lock an item  $X$  but is not able to because  $X$  is locked by some other transaction  $T_j$  with a conflicting lock.
- Wait-die. If  $TS(T_i) < TS(T_j)$ , then ( $T_i$  older than  $T_j$ )  $T_i$  is allowed to wait; otherwise ( $T_i$  younger than  $T_j$ ) abort  $T_i$  ( $T_i$  dies) and restart it later with the same timestamp.
- Wound-wait. If  $TS(T_i) < TS(T_j)$ , then ( $T_i$  older than  $T_j$ ) abort  $T_j$  ( $T_i$  wounds  $T_j$ ) and restart it later with the same timestamp; otherwise ( $T_i$  younger than  $T_j$ )  $T_i$  is allowed to wait.

# Database Concurrency Control

## Dealing with Deadlock and Starvation

### **Starvation**

Starvation occurs when a particular transaction consistently waits or restarted and never gets a chance to proceed further.

Solutions for starvation:

- To have a fair waiting scheme, such as using a first-come-first-served queue.
- To increase the priority of a transaction the longer it waits, until it eventually gets the highest priority and proceeds.
- To use higher priorities for transactions that have been aborted multiple times.
- The Wait-Die and Wound-Wait schemes avoid starvation, because they restart a transaction that has been aborted with its same original timestamp.



# Database Concurrency Control

## Timestamp based concurrency control algorithm

### ■ **Timestamp**

- A monotonically increasing variable (integer) indicating the age of an operation or a transaction. A larger timestamp value indicates a more recent event or operation.
- Timestamp based algorithm uses timestamp to serialize the execution of concurrent transactions.

# Database Concurrency Control

## Timestamp based concurrency control algorithm

### ■ Basic Timestamp Ordering

- 1. Transaction T issues a **write\_item(X)** operation:
  - If  $\text{read\_TS}(X) > \text{TS}(T)$  or if  $\text{write\_TS}(X) > \text{TS}(T)$ , then a younger transaction has already read the data item so abort and roll-back T and reject the operation.
  - If the condition in part (a) does not exist, then execute **write\_item(X)** of T and set  $\text{write\_TS}(X)$  to  $\text{TS}(T)$ .
- 2. Transaction T issues a **read\_item(X)** operation:
  - If  $\text{write\_TS}(X) > \text{TS}(T)$ , then a younger transaction has already written to the data item so abort and roll-back T and reject the operation.
  - If  $\text{write\_TS}(X) \leq \text{TS}(T)$ , then execute **read\_item(X)** of T and set  $\text{read\_TS}(X)$  to the larger of  $\text{TS}(T)$  and the current  $\text{read\_TS}(X)$ .

# Database Concurrency Control

## Timestamp based concurrency control algorithm

### ■ **Strict Timestamp Ordering**

- 1. Transaction T issues a `write_item(X)` operation:
  - If  $TS(T) > read\_TS(X)$ , then delay T until the transaction T' that wrote or read X has terminated (committed or aborted).
- 2. Transaction T issues a `read_item(X)` operation:
  - If  $TS(T) > write\_TS(X)$ , then delay T until the transaction T' that wrote or read X has terminated (committed or aborted).

# Database Concurrency Control

## Timestamp based concurrency control algorithm

### ■ **Thomas's Write Rule**

- If  $\text{read\_TS}(X) > \text{TS}(T)$  then abort and roll-back  $T$  and reject the operation.
- If  $\text{write\_TS}(X) > \text{TS}(T)$ , then just ignore the write operation and continue execution. This is because the most recent writes counts in case of two consecutive writes.
- If the conditions given in 1 and 2 above do not occur, then execute  $\text{write\_item}(X)$  of  $T$  and set  $\text{write\_TS}(X)$  to  $\text{TS}(T)$ .